

dog_app

April 2, 2020

1 Convolutional Neural Networks

1.1 Project: Write an Algorithm for a Dog Identification App

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with '**(IMPLEMENTATION)**' in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

Note: Once you have completed all of the code implementations, you need to finalize your work by exporting the Jupyter Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to **File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question X**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

Note: Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.

The rubric contains *optional* "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. If you decide to pursue the "Stand Out Suggestions", you should include the code in this Jupyter notebook.

Step 0: Import Datasets

Make sure that you've downloaded the required human and dog datasets:

Note: if you are using the Udacity workspace, you DO NOT need to re-download these - they can be found in the /data folder as noted in the cell below.

- Download the [dog dataset](#). Unzip the folder and place it in this project's home directory, at the location /dog_images.
- Download the [human dataset](#). Unzip the folder and place it in the home directory, at location /lfw.

Note: If you are using a Windows machine, you are encouraged to use [7zip](#) to extract the folder.

In the code cell below, we save the file paths for both the human (LFW) dataset and dog dataset in the numpy arrays human_files and dog_files.

```
In [1]: import numpy as np
        from glob import glob

        # load filenames for human and dog images
        human_files = np.array(glob("/data/lfw/*/.*"))
        dog_files = np.array(glob("/data/dog_images/*/.*"))

        # print number of images in each dataset
        print('There are %d total human images.' % len(human_files))
        print('There are %d total dog images.' % len(dog_files))
```

There are 13233 total human images.

There are 8351 total dog images.

Step 1: Detect Humans

In this section, we use OpenCV's implementation of [Haar feature-based cascade classifiers](#) to detect human faces in images.

OpenCV provides many pre-trained face detectors, stored as XML files on [github](#). We have downloaded one of these detectors and stored it in the haarcascades directory. In the next code cell, we demonstrate how to use this detector to find human faces in a sample image.

```
In [2]: import cv2
        import matplotlib.pyplot as plt
        %matplotlib inline

        # extract pre-trained face detector
        face_cascade = cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_alt.xml')

        # load color (BGR) image
        img = cv2.imread(human_files[0])
        # convert BGR image to grayscale
        gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

        # find faces in image
        faces = face_cascade.detectMultiScale(gray)

        # print number of faces detected in the image
        print('Number of faces detected:', len(faces))
```

```

# get bounding box for each detected face
for (x,y,w,h) in faces:
    # add bounding box to color image
    cv2.rectangle(img,(x,y),(x+w,y+h),(255,0,0),2)

# convert BGR image to RGB for plotting
cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

# display the image, along with bounding box
plt.imshow(cv_rgb)
plt.show()

```

Number of faces detected: 1



Before using any of the face detectors, it is standard procedure to convert the images to grayscale. The `detectMultiScale` function executes the classifier stored in `face_cascade` and takes the grayscale image as a parameter.

In the above code, `faces` is a numpy array of detected faces, where each row corresponds to a detected face. Each detected face is a 1D array with four entries that specifies the bounding box of the detected face. The first two entries in the array (extracted in the above code as `x` and `y`) specify the horizontal and vertical positions of the top left corner of the bounding box. The last two entries in the array (extracted here as `w` and `h`) specify the width and height of the box.

1.1.1 Write a Human Face Detector

We can use this procedure to write a function that returns True if a human face is detected in an image and False otherwise. This function, aptly named `face_detector`, takes a string-valued file path to an image as input and appears in the code block below.

```
In [3]: # returns "True" if face is detected in image stored at img_path
def face_detector(img_path):
    img = cv2.imread(img_path)
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    faces = face_cascade.detectMultiScale(gray)
    return len(faces) > 0
```

1.1.2 (IMPLEMENTATION) Assess the Human Face Detector

Question 1: Use the code cell below to test the performance of the `face_detector` function.

- What percentage of the first 100 images in `human_files` have a detected human face?
- What percentage of the first 100 images in `dog_files` have a detected human face?

Ideally, we would like 100% of human images with a detected face and 0% of dog images with a detected face. You will see that our algorithm falls short of this goal, but still gives acceptable performance. We extract the file paths for the first 100 images from each of the datasets and store them in the numpy arrays `human_files_short` and `dog_files_short`.

Answer: (You can print out your results and/or write your percentages in this cell) 98%, 17%

```
In [7]: from tqdm import tqdm
```

```
human_files_short = human_files[:100]
dog_files_short = dog_files[:100]
```

```
##-## Do NOT modify the code above this line. ##-##
```

```
## TODO: Test the performance of the face_detector algorithm
## on the images in human_files_short and dog_files_short.
```

```
human_count = 0
```

```
dog_count = 0
```

```
for image in human_files_short:
    if face_detector(image) == True:
        human_count += 1
```

```
for image in dog_files_short:
    if face_detector(image) == True:
        dog_count += 1
```

```
print(f'Percentage of human images that detected human face: {human_count}%')
```

```
print(f'Percentage of misclassified dog images that detected human faces: {dog_count}%')
```

Percentage of human images that detected human face: 98%

Percentage of misclassified dog images that detected human faces: 17%

We suggest the face detector from OpenCV as a potential way to detect human images in your algorithm, but you are free to explore other approaches, especially approaches that make use of deep learning :). Please use the code cell below to design and test your own face detection algorithm. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [8]: ### (Optional)  
       ### TODO: Test performance of another face detection algorithm.  
       ### Feel free to use as many code cells as needed.
```

Step 2: Detect Dogs

In this section, we use a [pre-trained model](#) to detect dogs in images.

1.1.3 Obtain Pre-trained VGG-16 Model

The code cell below downloads the VGG-16 model, along with weights that have been trained on [ImageNet](#), a very large, very popular dataset used for image classification and other vision tasks. ImageNet contains over 10 million URLs, each linking to an image containing an object from one of [1000 categories](#).

```
In [9]: import torch  
       import torchvision.models as models  
  
       # define VGG16 model  
       VGG16 = models.vgg16(pretrained=True)  
  
       # check if CUDA is available  
       use_cuda = torch.cuda.is_available()  
  
       # move model to GPU if CUDA is available  
       if use_cuda:  
           VGG16 = VGG16.cuda()
```

Given an image, this pre-trained VGG-16 model returns a prediction (derived from the 1000 possible categories in ImageNet) for the object that is contained in the image.

1.1.4 (IMPLEMENTATION) Making Predictions with a Pre-trained Model

In the next code cell, you will write a function that accepts a path to an image (such as `'dogImages/train/001.Affenpinscher/Affenpinscher_00001.jpg'`) as input and returns the index corresponding to the ImageNet class that is predicted by the pre-trained VGG-16 model. The output should always be an integer between 0 and 999, inclusive.

Before writing the function, make sure that you take the time to learn how to appropriately pre-process tensors for pre-trained models in the [PyTorch documentation](#).

```
In [10]: from PIL import Image  
        import torchvision.transforms as transforms
```

```

def load_image(img_path):
    ''' Load in and transform an image'''
    image = Image.open(img_path).convert('RGB')

    # VGG-16 Takes 224x224 images as input, so we resize all of them and convert data
    in_transform = transforms.Compose([
        transforms.Resize(224),
        transforms.CenterCrop(224),
        transforms.ToTensor(),
        transforms.Normalize((0.485, 0.456, 0.406),
                              (0.229, 0.224, 0.225))])

    # discard the transparent, alpha channel (that's the :3) and add the batch dimension
    image = in_transform(image)[:3,:,:].unsqueeze(0)

    return image

def VGG16_predict(img_path):
    '''
    Use pre-trained VGG-16 model to obtain index corresponding to
    predicted ImageNet class for image at specified path

    Args:
    img_path: path to an image

    Returns:
    Index corresponding to VGG-16 model's prediction
    '''

    ## TODO: Complete the function.
    ## Load and pre-process an image from the given img_path
    ## Return the *index* of the predicted class for that image

    image = load_image(img_path)
    if use_cuda:
        image = image.cuda()
    predict = VGG16(image)
    predict = predict.data.cpu().argmax()

    return predict # predicted class index

```

1.1.5 (IMPLEMENTATION) Write a Dog Detector

While looking at the [dictionary](#), you will notice that the categories corresponding to dogs appear in an uninterrupted sequence and correspond to dictionary keys 151-268, inclusive, to include all categories from 'Chihuahua' to 'Mexican hairless'. Thus, in order to check to see if an image is

predicted to contain a dog by the pre-trained VGG-16 model, we need only check if the pre-trained model predicts an index between 151 and 268 (inclusive).

Use these ideas to complete the `dog_detector` function below, which returns True if a dog is detected in an image (and False if not).

```
In [11]: ### returns "True" if a dog is detected in the image stored at img_path
def dog_detector(img_path):
    ## TODO: Complete the function.
    class_index = VGG16_predict(img_path)
    return ((class_index >= 151) & ( class_index <= 268))  # true/false
```

1.1.6 (IMPLEMENTATION) Assess the Dog Detector

Question 2: Use the code cell below to test the performance of your `dog_detector` function.

- What percentage of the images in `human_files_short` have a detected dog?
- What percentage of the images in `dog_files_short` have a detected dog?

Answer: 1%, 100%

```
In [12]: ### TODO: Test the performance of the dog_detector function
### on the images in human_files_short and dog_files_short.
from tqdm import tqdm

# Initiallizing:

dog_percentage_in_human = 0
dog_percentage_in_dog = 0

for i in tqdm(range(len(human_files_short))):
    dog_percentage_in_human += int (dog_detector(human_files_short[i]))

for i in tqdm(range(len(dog_files_short))):
    dog_percentage_in_dog += int (dog_detector(dog_files_short[i]))

print (' Dog Percentage in Human dataset: {}% \n Dog Percentage in Dog Dataset: {}%'.fo

100%|| 100/100 [00:03<00:00, 30.42it/s]
100%|| 100/100 [00:04<00:00, 25.52it/s]

Dog Percentage in Human dataset: 1%
Dog Percentage in Dog Dataset: 100%
```

We suggest VGG-16 as a potential network to detect dog images in your algorithm, but you are free to explore other pre-trained networks (such as [Inception-v3](#), [ResNet-50](#), etc). Please use the code cell below to test other pre-trained PyTorch models. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [13]: ### (Optional)
        ### TODO: Report the performance of another pre-trained network.
        ### Feel free to use as many code cells as needed.
```

Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

Now that we have functions for detecting humans and dogs in images, we need a way to predict breed from images. In this step, you will create a CNN that classifies dog breeds. You must create your CNN *from scratch* (so, you can't use transfer learning *yet!*), and you must attain a test accuracy of at least 10%. In Step 4 of this notebook, you will have the opportunity to use transfer learning to create a CNN that attains greatly improved accuracy.

We mention that the task of assigning breed to dogs from images is considered exceptionally challenging. To see why, consider that *even a human* would have trouble distinguishing between a Brittany and a Welsh Springer Spaniel.

Brittany	Welsh Springer Spaniel
----------	------------------------

It is not difficult to find other dog breed pairs with minimal inter-class variation (for instance, Curly-Coated Retrievers and American Water Spaniels).

Curly-Coated Retriever	American Water Spaniel
------------------------	------------------------

Likewise, recall that labradors come in yellow, chocolate, and black. Your vision-based algorithm will have to conquer this high intra-class variation to determine how to classify all of these different shades as the same breed.

Yellow Labrador	Chocolate Labrador
-----------------	--------------------

We also mention that random chance presents an exceptionally low bar: setting aside the fact that the classes are slightly imbalanced, a random guess will provide a correct answer roughly 1 in 133 times, which corresponds to an accuracy of less than 1%.

Remember that the practice is far ahead of the theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

1.1.7 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at `dog_images/train`, `dog_images/valid`, and `dog_images/test`, respectively). You may find [this documentation on custom datasets](#) to be a useful resource. If you are interested in augmenting your training and/or validation data, check out the wide variety of [transforms](#)!

```
In [14]: import os
         from torchvision import datasets
         from PIL import ImageFile
         ImageFile.LOAD_TRUNCATED_IMAGES = True
         ### TODO: Write data loaders for training, validation, and test sets
         ## Specify appropriate transforms, and batch_sizes

         # number of subprocesses to use for data loading
         num_workers = 0
         # how many samples per batch to load
         batch_size = 20

         data_dir = '/data/dog_images/'
         train_dir = os.path.join(data_dir, 'train/')
         test_dir = os.path.join(data_dir, 'test/')
         valid_dir = os.path.join(data_dir, 'valid/')

         transform = { 'train':

         data_transforms = {
             'train' : transforms.Compose([transforms.Resize(224),
                                           transforms.CenterCrop(224),
                                           transforms.RandomHorizontalFlip(), # randomly flip an
                                           transforms.RandomRotation(10),
                                           transforms.ToTensor(),
                                           transforms.Normalize(mean=[0.485, 0.456, 0.406],
                                                                std=[0.229, 0.224, 0.225]))),

             'valid' : transforms.Compose([transforms.Resize(224),
                                           transforms.CenterCrop(224),
                                           transforms.ToTensor(),
                                           transforms.Normalize(mean=[0.485, 0.456, 0.406],
                                                                std=[0.229, 0.224, 0.225]))),

             'test' : transforms.Compose([transforms.Resize(224),
                                           transforms.CenterCrop(224),
                                           transforms.ToTensor(),
                                           transforms.Normalize(mean=[0.485, 0.456, 0.406],
                                                                std=[0.229, 0.224, 0.225]))])}
```

```

image_datasets = {'train' : datasets.ImageFolder(root=train_dir,transform=data_transforms),
                  'valid' : datasets.ImageFolder(root=valid_dir,transform=data_transforms),
                  'test' : datasets.ImageFolder(root=test_dir,transform=data_transforms)}

loaders_scratch = {'train' : torch.utils.data.DataLoader(image_datasets['train'],batch_size=batch_size),
                   'valid' : torch.utils.data.DataLoader(image_datasets['valid'],batch_size=batch_size),
                   'test' : torch.utils.data.DataLoader(image_datasets['test'],batch_size=batch_size)}

dataset_sizes = {x: len(image_datasets[x]) for x in ['train', 'valid', 'test']}
class_names = image_datasets['train'].classes
n_classes = len(class_names)

print( len(class_names))
print( len(image_datasets['train']))

print( len(image_datasets['valid']))

print( len(image_datasets['test']))

```

133
6680
835
836

Question 3: Describe your chosen procedure for preprocessing the data. - How does your code resize the images (by cropping, stretching, etc)? What size did you pick for the input tensor, and why? - Did you decide to augment the dataset? If so, how (through translations, flips, rotations, etc)? If not, why not?

Answer: 1. I resized the images to 244X244 using transforms.Resize(244) to avoid the bias, the long time, and to save memory size in case of the big images size. 2. Yes, as the dataset size is small so I decided to augment the dataset through randomly flip using transforms.RandomHorizontalFlip() and images rotation using transforms.RandomRotation(10) functions.

1.1.8 (IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. Use the template in the code cell below.

```

In [15]: import torch.nn as nn
          import torch.nn.functional as F

          # define the CNN architecture
          class Net(nn.Module):
              ### TODO: choose an architecture, and complete the class

```

```

def __init__(self):
    super(Net, self).__init__()
    ## Define layers of a CNN

    # convolutional layer (sees 32x32x3 image tensor)
    self.conv1 = nn.Conv2d(3, 16, 3, padding=1)
    # convolutional layer (sees 16x16x16 tensor)
    self.conv2 = nn.Conv2d(16, 32, 3, padding=1)
    # convolutional layer (sees 8x8x32 tensor)
    self.conv3 = nn.Conv2d(32, 64, 3, padding=1)
    # max pooling layer
    self.pool = nn.MaxPool2d(2, 2)
    # linear layer (64 * 28 * 28 -> 500)
    self.fc1 = nn.Linear(64 * 28 * 28, 500)
    # linear layer (500 -> 133)
    self.fc2 = nn.Linear(500, 133)
    # dropout layer (p=0.25)
    self.dropout = nn.Dropout(0.25)

def forward(self, x):
    ## Define forward behavior

    # add sequence of convolutional and max pooling layers
    x = self.pool(F.relu(self.conv1(x)))
    x = self.pool(F.relu(self.conv2(x)))
    x = self.pool(F.relu(self.conv3(x)))
    # flatten image input
    x = x.view(-1, 64 * 28 * 28)
    # add dropout layer
    x = self.dropout(x)
    # add 1st hidden layer, with relu activation function
    x = F.relu(self.fc1(x))
    # add dropout layer
    x = self.dropout(x)
    # add 2nd hidden layer, with relu activation function
    x = self.fc2(x)

    return x

### You so NOT have to modify the code below this line. ###

# instantiate the CNN
model_scratch = Net()

# move tensors to GPU if CUDA is available
if use_cuda:
    model_scratch.cuda()

```

```
In [16]: print(model_scratch)
```

```
Net(
  (conv1): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (conv2): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (conv3): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (fc1): Linear(in_features=50176, out_features=500, bias=True)
  (fc2): Linear(in_features=500, out_features=133, bias=True)
  (dropout): Dropout(p=0.25)
)
```

Question 4: Outline the steps you took to get to your final CNN architecture and your reasoning at each step.

Answer: 1. As the input is RGB image cropped into 224x224 pixel and the depth is 3 (3 colors), we will have our input with the shape of 3x224x224. 2. The desired no. of output is the same as no. of classes (here we have 133 classes in our images datasets). 3. I created 3 convolutional layers with relu activation function and one max pooling layer(2,2), the first layer takes (3,224,224) input and converted it into a depth of 16 layers, the filter used was 3x3 with stride of 1 and padding of 1 in order to not get rid of edges: (i)first convolutional layer (sees 32x32x3 image tensor). (ii)second convolutional layer (sees 16x16x16 tensor). (iii)third convolutional layer (sees 8x8x32 tensor). (iv)one pool layer (2,2) was used in order to reduce the size of the images to half. 4. Then flatten image input using view function to reshape the tensor. 5. Then, two full connected Linear layers with relu activation function were added, and dropout layers for the hidden layers with a percentage of 25% to avoid the bias: (i)first linear layer (64 28 28 -> 500). (ii)second linear layer (500 -> 133), as the output classes is 133.

1.1.9 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_scratch`, and the optimizer as `optimizer_scratch` below.

```
In [17]: import torch.optim as optim

        ### TODO: select loss function
        criterion_scratch = nn.CrossEntropyLoss()

        ### TODO: select optimizer
        optimizer_scratch = optim.SGD(model_scratch.parameters(), lr=0.01)
```

1.1.10 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_scratch.pt'`.

```
In [19]: def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
        """returns trained model"""
        # initialize tracker for minimum validation loss
```

```

valid_loss_min = np.Inf

for epoch in range(1, n_epochs+1):
    # initialize variables to monitor training and validation loss
    train_loss = 0.0
    valid_loss = 0.0

    #####
    # train the model #
    #####
    model.train()
    for batch_idx, (data, target) in enumerate(loaders['train']):
        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()
        ## find the loss and update the model parameters accordingly
        ## record the average training loss, using something like
        ## train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))

        # clear the gradients of all optimized variables
        optimizer.zero_grad()
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model(data)
        # calculate the batch loss
        loss = criterion(output, target)
        # backward pass: compute gradient of the loss with respect to model parameters
        loss.backward()
        # perform a single optimization step (parameter update)
        optimizer.step()
        # update training loss
        train_loss += loss.item()*data.size(0)

    #####
    # validate the model #
    #####
    model.eval()
    for batch_idx, (data, target) in enumerate(loaders['valid']):
        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()
        ## update the average validation loss
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model(data)
        # calculate the batch loss
        loss = criterion(output, target)
        # update average validation loss
        valid_loss += loss.item()*data.size(0)

```

```

train_loss = train_loss/len(loaders['train'].dataset) ###
valid_loss = valid_loss/len(loaders['valid'].dataset) ###

# print training/validation statistics
print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
    epoch,
    train_loss,
    valid_loss
))

## TODO: save the model if validation loss has decreased
if valid_loss <= valid_loss_min:
    print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'.format(
        valid_loss_min,
        valid_loss))
    torch.save(model.state_dict(), save_path)
    valid_loss_min = valid_loss

# return trained model
return model

# train the model
model_scratch = train(20, loaders_scratch, model_scratch, optimizer_scratch,
    criterion_scratch, use_cuda, 'model_scratch.pt')

# load the model that got the best validation accuracy
model_scratch.load_state_dict(torch.load('model_scratch.pt'))

```

```

Epoch: 1      Training Loss: 4.207766      Validation Loss: 4.303818
Validation loss decreased (inf --> 4.303818). Saving model ...
Epoch: 2      Training Loss: 4.121917      Validation Loss: 4.277724
Validation loss decreased (4.303818 --> 4.277724). Saving model ...
Epoch: 3      Training Loss: 4.038090      Validation Loss: 4.245886
Validation loss decreased (4.277724 --> 4.245886). Saving model ...
Epoch: 4      Training Loss: 3.937896      Validation Loss: 4.206313
Validation loss decreased (4.245886 --> 4.206313). Saving model ...
Epoch: 5      Training Loss: 3.857815      Validation Loss: 4.151789
Validation loss decreased (4.206313 --> 4.151789). Saving model ...
Epoch: 6      Training Loss: 3.770598      Validation Loss: 4.145334
Validation loss decreased (4.151789 --> 4.145334). Saving model ...
Epoch: 7      Training Loss: 3.663791      Validation Loss: 4.180049
Epoch: 8      Training Loss: 3.543664      Validation Loss: 4.078306
Validation loss decreased (4.145334 --> 4.078306). Saving model ...
Epoch: 9      Training Loss: 3.423640      Validation Loss: 4.163184
Epoch: 10     Training Loss: 3.305253      Validation Loss: 4.175038
Epoch: 11     Training Loss: 3.159356      Validation Loss: 4.171958

```

Epoch: 12	Training Loss: 3.011615	Validation Loss: 4.259992
Epoch: 13	Training Loss: 2.852608	Validation Loss: 4.310714
Epoch: 14	Training Loss: 2.703452	Validation Loss: 4.418943
Epoch: 15	Training Loss: 2.489482	Validation Loss: 4.423641
Epoch: 16	Training Loss: 2.315187	Validation Loss: 4.470788
Epoch: 17	Training Loss: 2.129676	Validation Loss: 4.640040
Epoch: 18	Training Loss: 1.981928	Validation Loss: 4.717066
Epoch: 19	Training Loss: 1.774967	Validation Loss: 4.822499
Epoch: 20	Training Loss: 1.651740	Validation Loss: 4.891258

1.1.11 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 10%.

```
In [20]: def test(loaders, model, criterion, use_cuda):

    # monitor test loss and accuracy
    test_loss = 0.
    correct = 0.
    total = 0.

    model.eval()
    for batch_idx, (data, target) in enumerate(loaders['test']):
        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model(data)
        # calculate the loss
        loss = criterion(output, target)
        # update average test loss
        test_loss = test_loss + ((1 / (batch_idx + 1)) * (loss.data - test_loss))
        # convert output probabilities to predicted class
        pred = output.data.max(1, keepdim=True)[1]
        # compare predictions to true label
        correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy())
        total += data.size(0)

    print('Test Loss: {:.6f}\n'.format(test_loss))

    print('\nTest Accuracy: %2d%% (%2d/%2d)' % (
        100. * correct / total, correct, total))

    # call test function
    test(loaders_scratch, model_scratch, criterion_scratch, use_cuda)
```

Test Loss: 4.090515

Test Accuracy: 8% (74/836)

Step 4: Create a CNN to Classify Dog Breeds (using Transfer Learning)

You will now use transfer learning to create a CNN that can identify dog breed from images. Your CNN must attain at least 60% accuracy on the test set.

1.1.12 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at `dogImages/train`, `dogImages/valid`, and `dogImages/test`, respectively).

If you like, **you are welcome to use the same data loaders from the previous step**, when you created a CNN from scratch.

```
In [21]: ## TODO: Specify data loaders
        loaders_transfer = loaders_scratch
```

1.1.13 (IMPLEMENTATION) Model Architecture

Use transfer learning to create a CNN to classify dog breed. Use the code cell below, and save your initialized model as the variable `model_transfer`.

```
In [22]: import torchvision.models as models
        import torch.nn as nn

        ## TODO: Specify model architecture
        # Load the pretrained model from pytorch
        model_transfer = models.vgg16(pretrained=True)

        # Freeze training for all "features" layers
        for param in model_transfer.features.parameters():
            param.requires_grad = False

        n_inputs = model_transfer.classifier[6].in_features

        # add last linear layer (n_inputs -> 133 classes)
        # new layers automatically have requires_grad = True
        last_layer = nn.Linear(n_inputs, 133)

        model_transfer.classifier[6] = last_layer

        if use_cuda:
            model_transfer = model_transfer.cuda()
```


Question 5: Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

Answer: 1. Pre-trained modules are well trained to extract patterns and features very well. 2. I chose to load the VGG16 models as transfer learning. I checked the structure then freeze training for all "features" layers because if I changed it will take a long time. I changed the last linear layer output to be 133 classes to fit with our classification output. Then train the classifier to get a good result for our application.

1.1.14 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_transfer`, and the optimizer as `optimizer_transfer` below.

```
In [23]: criterion_transfer = nn.CrossEntropyLoss()
         optimizer_transfer = optim.SGD(model_transfer.classifier.parameters(), lr=0.001)
```

1.1.15 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_transfer.pt'`.

```
In [31]: # train the model
         model_transfer = train(10, loaders_transfer, model_transfer, optimizer_transfer, criterion_transfer)
         # train(n_epochs, loaders_transfer, model_transfer, optimizer_transfer, criterion_transfer)

         # load the model that got the best validation accuracy (uncomment the line below)
         model_transfer.load_state_dict(torch.load('model_transfer.pt'))
```

```
Epoch: 1      Training Loss: 0.444545      Validation Loss: 0.402775
Validation loss decreased (inf --> 0.402775). Saving model ...
Epoch: 2      Training Loss: 0.417981      Validation Loss: 0.398333
Validation loss decreased (0.402775 --> 0.398333). Saving model ...
Epoch: 3      Training Loss: 0.396322      Validation Loss: 0.397795
Validation loss decreased (0.398333 --> 0.397795). Saving model ...
Epoch: 4      Training Loss: 0.381947      Validation Loss: 0.385124
Validation loss decreased (0.397795 --> 0.385124). Saving model ...
Epoch: 5      Training Loss: 0.354258      Validation Loss: 0.380130
Validation loss decreased (0.385124 --> 0.380130). Saving model ...
Epoch: 6      Training Loss: 0.348644      Validation Loss: 0.378747
Validation loss decreased (0.380130 --> 0.378747). Saving model ...
Epoch: 7      Training Loss: 0.338631      Validation Loss: 0.372815
Validation loss decreased (0.378747 --> 0.372815). Saving model ...
Epoch: 8      Training Loss: 0.332489      Validation Loss: 0.363387
Validation loss decreased (0.372815 --> 0.363387). Saving model ...
Epoch: 9      Training Loss: 0.309466      Validation Loss: 0.375936
Epoch: 10     Training Loss: 0.310998      Validation Loss: 0.365197
```

1.1.16 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 60%.

```
In [32]: test(loaders_transfer, model_transfer, criterion_transfer, use_cuda)
```

Test Loss: 0.426434

Test Accuracy: 87% (730/836)

1.1.17 (IMPLEMENTATION) Predict Dog Breed with the Model

Write a function that takes an image path as input and returns the dog breed (Affenpinscher, Afghan hound, etc) that is predicted by your model.

```
In [33]: ### TODO: Write a function that takes a path to an image as input
         ### and returns the dog breed that is predicted by the model.

         # list of class names by index, i.e. a name can be accessed like class_names[0]
         class_names = [item[4:].replace("_", " ") for item in image_datasets['train'].classes]

def predict_breed_transfer(img_path):
    # load the image and return the predicted breed

    image = load_image(img_path)
    if use_cuda:
        image = image.cuda()
    predict = model_transfer(image)
    predict = predict.data.cpu().argmax()

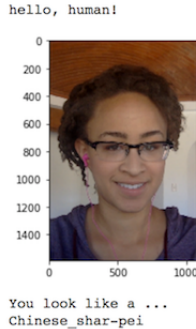
    return class_names[predict]
```

Step 5: Write your Algorithm

Write an algorithm that accepts a file path to an image and first determines whether the image contains a human, dog, or neither. Then, - if a **dog** is detected in the image, return the predicted breed. - if a **human** is detected in the image, return the resembling dog breed. - if **neither** is detected in the image, provide output that indicates an error.

You are welcome to write your own functions for detecting humans and dogs in images, but feel free to use the `face_detector` and `human_detector` functions developed above. You are **required** to use your CNN from Step 4 to predict dog breed.

Some sample output for our algorithm is provided below, but feel free to design your own user experience!



Sample Human Output

1.1.18 (IMPLEMENTATION) Write your Algorithm

In [34]: *### TODO: Write your algorithm.*

Feel free to use as many code cells as needed.

```
def display_img(img_path):
    img = Image.open(img_path)
    plt.imshow(img)
    plt.show()

def run_app(img_path):
    ## handle cases for a human face, dog, and neither

    predicted_breed = predict_breed_transfer (img_path)
    if dog_detector(img_path):
        print ('\n\n hello, it is a dog ')
        display_img(img_path)
        return print ('the predicted breed is:', predicted_breed)

    elif face_detector(img_path):
        print ('\n\n hello, human')
        display_img(img_path)
        return print (' You look like a ', predicted_breed)

    else:
        display_img(img_path)
        print ('\n\n error: sorry it is neither human, nor dog ')
    predict_breed_transfer
```

Step 6: Test Your Algorithm

In this section, you will take your new algorithm for a spin! What kind of dog does the algorithm think that *you* look like? If you have a dog, does it predict your dog's breed accurately? If

you have a cat, does it mistakenly think that your cat is a dog?

1.1.19 (IMPLEMENTATION) Test Your Algorithm on Sample Images!

Test your algorithm at least six images on your computer. Feel free to use any images you like. Use at least two human and two dog images.


Question 6: Is the output better than you expected :) ? Or worse :(? Provide at least three possible points of improvement for your algorithm.

Answer: (Three possible points for improvement) 1. If we trained the model with a dataset has bigger number of images it will make an improvement. 2. If we consider more images augmentation this will make an improvement. 3. If we trained the model with more number of epochs this could make an improvement.

```
In [35]: ## TODO: Execute your algorithm from Step 6 on
        ## at least 6 images on your computer.
        ## Feel free to use as many code cells as needed.


        ## suggested code, below
        for file in np.hstack((human_files[:3], dog_files[:3])):
            run_app(file)
```

hello, human




You look like a American foxhound

hello, human




You look like a Dachshund

hello, human




You look like a American water spaniel

hello, it is a dog




the predicted breed is: Bullmastiff

hello, it is a dog



the predicted breed is: Mastiff

hello, it is a dog



the predicted breed is: Bullmastiff